

Combinatorial Pure Exploration of Multi-Armed Bandits

Shouyuan Chen¹ Tian Lin² Irwin King¹ Michael R. Lyu¹ Wei
Chen³


¹ CUHK ² Tsinghua University ³ Microsoft Research Asia

Single-armed bandit



arm
(slot machine)

Single-armed bandit




Single-armed bandit



Multi-armed bandit


n arms



1. each arm has an **unknown** reward distribution
2. the reward distributions can be **different**.

Multi-armed bandit

n arms



a game on multiple rounds...


rules

for round $t = 1, \dots, T$

- plays arm $i_t \in [n]$
- receives reward $X_{it} \sim \phi_i$

Multi-armed bandit

n arms



a game on multiple rounds...


rules

for round $t = 1, \dots, T$

- plays arm $i_t \in [n]$
- receives reward $X_{it} \sim \phi_i$

Multi-armed bandit

n arms



a game on multiple rounds...


rules

for round $t = 1, \dots, T$

- plays arm $i_t \in [n]$
- receives reward $X_{it} \sim \phi_i$

Multi-armed bandit

n arms



a game on multiple rounds...


rules

for round $t = 1, \dots, T$

- plays arm $i_t \in [n]$
- receives reward $X_{it} \sim \phi_i$

Multi-armed bandit

n arms




rules

for round $t = 1, \dots, T$

- plays arm $i_t \in [n]$
- receives reward $X_{it} \sim \phi_i$

Multi-armed bandit

n arms



rules

for round $t = 1, \dots, T$

- plays arm $i_t \in [n]$
- receives reward $X_{it} \sim \phi_i$

in the end...

take all rewards




goal: maximize the cumulative reward


exploitation v.s. exploration

Multi-armed bandit



n arms



n arms pure exploration



in the end...

take all rewards   

goal: maximize the cumulative reward


exploitation v.s. exploration

Multi-armed bandit

rules

for round $t = 1, \dots, T$

- plays arm $i_t \in [n]$
- receives reward $X_{it} \sim \phi_i$




Multi-armed bandit


rules

for round $t = 1, \dots, T$

- plays arm $i_t \in [n]$
- receives reward $X_{it} \sim \phi_i$



in the end...


- (1) forfeit all rewards
 - (2) output **1** arm
- 

goal: find the single **best arm**
(largest expected reward)



Pure Exploration of MAB

A/B testing, clinical trials, wireless network, crowdsourcing, ...



- n arms = n variants
- play arm i = a page view on the i -th variant
- reward = a click on the ads
- finding the best arm = finding the variant with the highest average ads clicks

Pure exploration: two settings

fixed budget

- play for T rounds.
- report the best arm after finished.
- **goal:** minimize the probability of error $\Pr[\text{out} \neq i_*]$

fixed confidence

- play for any number of rounds.
- report the best arm after finished
- guarantee that probability of error $\Pr[\text{out} \neq i_*] < \delta$.
- **goal:** minimize the number of rounds (**sample complexity**).

Combinatorial Pure Exploration of MAB

Combinatorial Pure Exploration (CPE)

- play one arm at each round
- find the optimal **set** of arms M_* satisfying certain constraint


$$M_* = \arg \max_{M \in \mathcal{M}} \sum_{e \in M} w(e)$$

- ▶ $[n]$: set of arms
- ▶ $\mathcal{M} \subseteq 2^{[n]}$: **decision class** with a combinatorial constraint
- ▶ maximize the **sum of expected rewards** of arms in the set


size- k -sets



spanning trees



paths




matchings



Motivating Examples

- matching




Goal:

- 1) estimate the productivities from tests.
- 2) find the optimal **1-1 assignment**.

Motivating Examples


- matching



Goal:

- 1) estimate the productivities from tests.
- 2) find the optimal **1-1 assignment**.

- spanning trees and paths




Goal:

- 1) estimate the delays from measurements
- 2) find the **minimum spanning tree** or **shortest path**.

Motivating Examples


- matching



Goal:

- 1) estimate the productivities from tests.
- 2) find the optimal **1-1 assignment**.

- spanning trees and paths



Goal:

- 1) estimate the delays from measurements
- 2) find the **minimum spanning tree** or **shortest path**.

- size- k -sets

- ▶ finding the top- k arms.

Existing Work

- find top- k arms [KS10,GGL12,KTPS12,BWV13,KK13,ZCL14]
- find top arms in disjoint groups of arms
[GGLB11,GGL12,BWV13]
- separate treatments, no unified framework

Our Results

Our Results

- general framework
 - ▶ for a wide range of combinatorial constraints \mathcal{M} .

Our Results

- **general framework**
 - ▶ for a wide range of combinatorial constraints \mathcal{M} .
- **algorithms**
 - ▶ two generic learning algorithms.

Our Results

- **general framework**
 - ▶ for a wide range of combinatorial constraints \mathcal{M} .
- **algorithms**
 - ▶ two generic learning algorithms.
- **upper bounds**
 - ▶ sample complexity / probability of error.

Our Results

- **general framework**
 - ▶ for a wide range of combinatorial constraints \mathcal{M} .
- **algorithms**
 - ▶ two generic learning algorithms.
- **upper bounds**
 - ▶ sample complexity / probability of error.
- **lower bound**
 - ▶ algorithms are **optimal** (within log factors) for many types of \mathcal{M} (in particular, bases of a matroid).

Our Results

- **general framework**
 - ▶ for a wide range of combinatorial constraints \mathcal{M} .
- **algorithms**
 - ▶ two generic learning algorithms.
- **upper bounds**
 - ▶ sample complexity / probability of error.
- **lower bound**
 - ▶ algorithms are **optimal** (within log factors) for many types of \mathcal{M} (in particular, bases of a matroid).
- **compared with existing work**
 - ▶ the first lower bound for the top- k problem
 - ▶ the first upper and lower bounds for other combinatorial constraints.

CLUCB: Fixed confidence algorithm


input

- confidence: $\delta \in (0, 1)$
- access to a maximization oracle: $\text{Oracle}(\cdot) : \mathbb{R}^n \rightarrow \mathcal{M}$
 - ▶ $\text{Oracle}(v) = \arg \max_{M \in \mathcal{M}} \sum_{i \in M} v(i)$ for weights $v \in \mathbb{R}^n$

output

- a set of arms: $M \in \mathcal{M}$.


CLUCB: Fixed confidence algorithm



notations

- for each arm $i \in [n]$ in each round t
 - ▶ empirical mean: $\bar{w}_t(i)$
 - ▶ confidence radius: $\text{rad}_t(i)$ (proportional to $1/\sqrt{n_t(i)}$)

CLUCB: Fixed confidence algorithm




CLUCB: Fixed confidence algorithm


all arms




maintain: for all i and t



Step 1



Step 2



$-\text{rad}_t(i)$

$+\text{rad}_t(i)$


$$\tilde{w}_t(i) = \bar{w}_t(i) \pm \text{rad}_t(i)$$

CLUCB: Fixed confidence algorithm


all arms




maintain: for all i and t



Step 1




Step 2



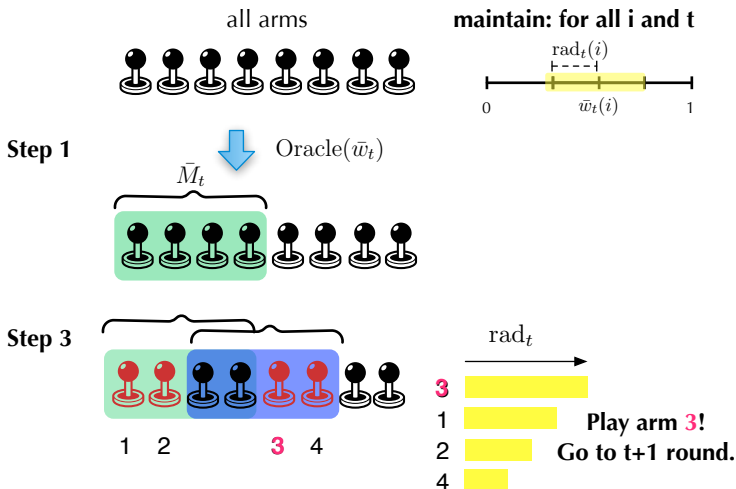
$$\tilde{w}_t(i) = \bar{w}_t(i) \pm \text{rad}_t(i)$$

If: $\bar{M}_t = \tilde{M}_t$
Then: Stop and output \bar{M}_t

CLUCB: Fixed confidence algorithm



CLUCB: Fixed confidence algorithm



CLUCB: Sample Complexity

Our sample complexity bound depends on two quantities.

- H : depends on expected rewards
- $\text{width}(\mathcal{M})$: depends on the structure of \mathcal{M}

CLUCB: Sample Complexity

Theorem (Upper bound)

With probability at least $1 - \delta$, CLUCB algorithm:

- 1. correctly outputs the optimal set M_**
- 2. uses at most $O(\text{width}(\mathcal{M})^2 \mathbf{H} \log(n\mathbf{H}/\delta))$ rounds.*

Our sample complexity bound depends on two quantities.

- \mathbf{H} : depends on expected rewards
- $\text{width}(\mathcal{M})$: depends on the structure of \mathcal{M}

Results at a glance

Theorem (Upper bound)

With probability at least $1 - \delta$, CLUCB algorithm:

1. outputs the optimal set $M_* \triangleq \arg \max_{M \in \mathcal{M}} w(M)$.
2. uses at most $O(\text{width}(\mathcal{M})^2 \mathbf{H} \log(n\mathbf{H}/\delta))$ rounds.

Theorem (Lower bound)

Given any expected rewards, any δ -correct algorithm must use at least $\Omega(\mathbf{H} \log(1/\delta))$ rounds. (An algorithm \mathbb{A} is δ -correct algorithm, if \mathbb{A} 's probability of error is at most δ for any instances)

Example (Sample Complexities)

- k -sets, spanning trees, bases of a matroid: $\tilde{O}(\mathbf{H})$ optimal!
- matchings, paths (in DAG): $\tilde{O}(|V|^2 \mathbf{H})$.
- in general: $\tilde{O}(n^2 \mathbf{H})$

H and gaps

- Δ_e : **gap** of arm $e \in [n]$

$$\Delta_e = \begin{cases} w(M_*) - \max_{M \in \mathcal{M}: e \in M} w(M) & \text{if } e \notin M_*, \\ w(M_*) - \max_{M \in \mathcal{M}: e \notin M} w(M) & \text{if } e \in M_* \end{cases}$$

- ▶ stability of the optimality of M_* wrt. arm e .
- $\mathbf{H} = \sum_{e \in [n]} \Delta_e^{-2}$
 - ▶ for the **top- K** problem: recover the previous definition of \mathbf{H} .

Width and exchange class


Intuitions

- we need a unifying method of analyzing different \mathcal{M}
 - ▶ an **exchange class** is a “proxy” for the structure of \mathcal{M} .
- an exchange class \mathcal{B} is a collection of “patches”
 $((b_+, b_-), b_+, b_- \subseteq [n])$ that are used to interpolate between valid sets $(M \setminus b_- \cup b_+ = M', M, M' \in \mathcal{M})$.


size- k -sets




spanning tree



matching



path




Width and exchange class

definition

$\text{width}(\mathcal{B})$: the size of the largest “patch”

$$\text{width}(\mathcal{B}) = \max_{(b_+, b_-) \in \mathcal{B}} |b_+| + |b_-|.$$



$\text{width}(\mathcal{M})$: the width of the “thinnest” exchange class

$$\text{width}(\mathcal{M}) = \min_{\mathcal{B} \in \text{Exchange}(\mathcal{M})} \text{width}(\mathcal{B}),$$

The main technical contribution: Define exchange class and its algebra and conduct generic analysis using exchange classes.

Example (Widths)

- **k -sets, spanning trees, bases of a matroid**: $\text{width}(\mathcal{M}) = 2$.
- **matchings, paths (in DAG)**: $\text{width}(\mathcal{M}) = O(|V|)$.
- **in general**: $\text{width}(\mathcal{M}) \leq n$

CSAR: Fixed budget algorithm

input

- budget: T (play for at most T rounds)
- access to a maximization oracle

output

- a set of arms: $M \in \mathcal{M}$.


overview:

- break the T rounds into n phases.


CSAR: Fixed budget algorithm

in each phase (n phases in total):


phase 1



phase 2



phase 3



- 1 arm is **accepted** or **rejected**.
- **active arms** are sampled for a same number of times.



active: neither accepted nor rejected.
require more samples



accepted: include in the output





rejected: exclude from the output

CSAR: Fixed budget algorithm


in each phase (n phases in total):

- 1 arm is **accepted** or **rejected**.
- active arms** are sampled for a same number of times.

 **active**: neither accepted nor rejected.
require more samples

 **accepted**: include in the output


 **rejected**: exclude from the output




CSAR: Fixed budget algorithm


in each phase (n phases in total):

- 1 arm is **accepted** or **rejected**.
- **active arms** are sampled for a same number of times.

 **active:** neither accepted nor rejected.
require more samples

 **accepted:** include in the output

 **rejected:** exclude from the output



problem: which arm to accept or reject?

CSAR: Fixed budget algorithm

problem: which arm to accept or reject?

- accept/reject the arm with the largest **empirical gap**.

$$\bar{\Delta}_e = \begin{cases} \bar{w}_t(\bar{M}_t) - \max_{M \in \mathcal{M}_t: e \in M} \bar{w}_t(M) & \text{if } e \notin \bar{M}_t, \\ \bar{w}_t(\bar{M}_t) - \max_{M \in \mathcal{M}_t: e \notin M} \bar{w}_t(M) & \text{if } e \in \bar{M}_t \end{cases}$$

- ▶ $\mathcal{M}_t = \{M : M \in \mathcal{M}, \textcolor{green}{A}_t \subseteq M, \textcolor{red}{B}_t \cap M = \emptyset\}$.
 - ▶ $\textcolor{green}{A}_t$: accepted arms, $\textcolor{red}{B}_t$: rejected arms (up to phase t).
 - ▶ $\rightarrow \bar{\Delta}_e$ can be computed using a maximization oracle.
- \rightarrow recall the (unknown) **gap** of arm e :

$$\Delta_e = \begin{cases} w(M_*) - \max_{M \in \mathcal{M}: e \in M} w(M) & \text{if } e \notin M_*, \\ w(M_*) - \max_{M \in \mathcal{M}: e \notin M} w(M) & \text{if } e \in M_* \end{cases}$$

CSAR: Probability of error

Theorem (Probability of error of CSAR)

Given any budget $T > n$, CSAR correctly outputs the optimal set M_ with probability at least*

$$1 - 2^{\tilde{O}\left(-\frac{T}{\text{width}(\mathcal{M})^2 \mathbf{H}}\right)}$$

and uses at most T rounds.

CSAR: Probability of error

Theorem (Probability of error of CSAR)

Given any budget $T > n$, CSAR correctly outputs the optimal set M_ with probability at least*

$$1 - 2^{\tilde{O}\left(-\frac{T}{\text{width}(\mathcal{M})^2 \mathbf{H}}\right)}$$

and uses at most T rounds.

Remark: To guarantee a constant probability of error of δ , both CSAR and CLUCB need $T = \tilde{O}(\text{width}(\mathcal{M})^2 \mathbf{H})$ rounds.

Summary

- **Combinatorial pure exploration**: a general framework that covers many pure exploration problems in MAB.
 - ▶ find top- k arms, optimal spanning trees, matchings or paths.
- Two general algorithms (CLUCB, CSAR) for the problem
 - ▶ only need a maximization oracle for \mathcal{M} .
 - ▶ comparable performance guarantees.
- Our algorithm is optimal (up to log factors) for matroids.
 - ▶ including k -sets and spanning trees.
- Trilogy on stochastic and combinatorial online learning : together with our recent work on **combinatorial multi-armed bandit** [CWY, ICML'13] and **combinatorial partial monitoring** [LAKLC, ICML'14], all dealing with general combinatorial constraints

Future work

- Tighten the bounds for matching, paths and other combinatorial constraints
- Support approximation oracles
- Support non-linear reward functions

Thank you!

See you at Poster Wed2 tonight!

Exchange class: Formal definition

Exchange set

An **exchange set** b is an ordered pair of disjoint sets $b = (b_+, b_-)$ where $b_+ \cap b_- = \emptyset$ and $b_+, b_- \subseteq [n]$.




Let M be any set. We also define two operators:

- $M \oplus b \triangleq M \setminus b_- \cup b_+$.
- $M \ominus b \triangleq M \setminus b_+ \cup b_-$.

Exchange class

We call a collection of exchange sets \mathcal{B} an **exchange class for \mathcal{M}** if \mathcal{B} satisfies the following property. For any $M, M' \in \mathcal{M}$ such that $M \neq M'$ and for any $e \in (M \setminus M')$, there exists an exchange set $(b_+, b_-) \in \mathcal{B}$ which satisfies five constraints: **(a)** $e \in b_-$, **(b)** $b_+ \subseteq M' \setminus M$, **(c)** $b_- \subseteq M \setminus M'$, **(d)** $(M \oplus b) \in \mathcal{M}$ and **(e)** $(M' \ominus b) \in \mathcal{M}$.

Experiments of CPE



Width and exchange class

definition

Let \mathcal{B} be an exchange class.

$$\text{width}(\mathcal{B}) = \max_{(b_+, b_-) \in \mathcal{B}} |b_+| + |b_-|.$$

Let $\text{Exchange}(\mathcal{M})$ denote the family of all possible exchange classes for \mathcal{M} . We define the width of \mathcal{M} to be the width of the thinnest exchange class

$$\text{width}(\mathcal{M}) = \min_{\mathcal{B} \in \text{Exchange}(\mathcal{M})} \text{width}(\mathcal{B}),$$

where $\text{Exchange}(\mathcal{M})$ is the family of all possible exchange classes for \mathcal{M} .